

# 02456 Deep Learning Written exam (solutions)

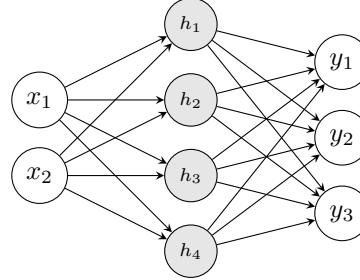
Technical University of Denmark

December 10, 2025

This document contains the 10 exam questions on 7 pages. **All answers must be filled in on the *answer sheet*; do not hand in this document.** Please read the instructions on the *answer sheet* carefully.

## Question 1

Consider the fully connected neural network shown below, where  $\mathbf{x} \in \mathbb{R}^2$  is the input and  $\mathbf{y} \in \mathbb{R}^3$  is the output.



Assume each neuron has a bias term. How many parameters does the model have in total?

- A. 7
- B. 9
- C. 20
- D. 27**
- E. Don't know

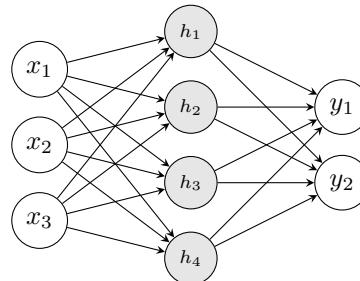
**Solution:** It has  $2 \times 4 + 4 \times 3 = 20$  weights and  $4 + 3 = 7$  biases, which is 27 in total.

## Question 2

Consider the fully connected neural network with one hidden layer defined as

$$f(\mathbf{x}) = W^{(2)}\sigma(W^{(1)}\mathbf{x}), \quad (1)$$

where  $W^{(1)} \in \mathbb{R}^{n \times m}$ ,  $W^{(2)} \in \mathbb{R}^{k \times \ell}$  and  $\sigma$  is the non-linear activation function. Assume that the network has the architecture shown below.



What are the dimensions of the weight matrices?

- A.  $n = 3, m = 4, k = 4, \ell = 2$
- B.  $n = 4, m = 3, k = 2, \ell = 4$
- C.  $n = 2, m = 4, k = 4, \ell = 3$
- D.  $n = 4, m = 2, k = 3, \ell = 4$
- E. Don't know

**Solution:** From the figure, we have that  $\dim(\mathbf{x}) = 3$ ,  $\dim(\mathbf{h}) = 4$  and  $\dim(\mathbf{y}) = 2$ . The matrix  $W^{(1)}$ , must therefore have dimensions  $\dim(\mathbf{h}) \times \dim(\mathbf{x}) = 4 \times 3$  and accordingly  $n = 4$  and  $m = 3$ . Similarly,  $W^{(2)}$  must have dimensions  $\dim(\mathbf{y}) \times \dim(\mathbf{h}) = 2 \times 4$  and accordingly  $k = 2$  and  $\ell = 4$ .

### Question 3

Consider the following simple neural network with a single hidden layer:

$$h = \text{ReLU}(wx), \quad y = vh, \quad L = \frac{1}{2}(y - t)^2, \quad (2)$$

where  $x \in \mathbb{R}$  is the input,  $h \in \mathbb{R}$  is the hidden layer output,  $y \in \mathbb{R}$  is the output,  $L \in \mathbb{R}$  is the loss,  $t \in \mathbb{R}$  is the target,  $w, v \in \mathbb{R}$  are weights, and  $\text{ReLU}(z) = \max(0, z)$ . Assume that  $x = 1$ ,  $t = 6$ ,  $w = 2$ , and  $v = 4$ . What is the value of the gradient of the loss  $\frac{\partial L}{\partial w}$  for this input and target?

- A. 0
- B. 4
- C. 8
- D. 16
- E. Don't know

**Solution:** First, we evaluate the forward pass:

$$h = \text{ReLU}(wx) = \text{ReLU}(2) = 2, \quad y = vh = 8, \quad L = \frac{1}{2}(8 - 6)^2 = 2. \quad (S1)$$

Using the chain rule, we find that

$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial y} \frac{\partial y}{\partial h} \frac{\partial h}{\partial w}, \quad (S2)$$

and compute each term as

$$\frac{\partial L}{\partial y} = y - t = 2, \quad \frac{\partial y}{\partial h} = v = 4, \quad \frac{\partial h}{\partial w} = \text{ReLU}'(wx) \cdot \frac{\partial(wx)}{\partial w} = 1 \cdot x = 1, \quad (S3)$$

where we used that

$$\text{ReLU}'(z) = \begin{cases} 1 & z > 0, \\ 0 & z < 0. \end{cases} \quad (S4)$$

Therefore, we find that

$$\frac{\partial L}{\partial w} = 2 \cdot 4 \cdot 1 = 8. \quad (S5)$$

### Question 4

Consider constructing a fully connected network with two hidden layers of size 128 that takes a  $10 \times 10$  image with three channels as input and classifies it into five classes. All internal nodes use the ReLU activation function. The input tensor has shape `[B, 3, 10, 10]`, where `B` is the batch size, and `nn` has been imported from `torch`. Which one of the following PyTorch code snippets implements such a network correctly?

A.

```

1 nn.Sequential(
2     nn.Linear(3, 10, 10, 128),
3     nn.ReLU(),
4     nn.Linear(128, 128),
5     nn.ReLU(),
6     nn.Linear(128, 5),
7     nn.Softmax(dim=1)
8 )

```

B.

```

1 nn.Sequential(
2     nn.Flatten(128),
3     nn.ReLU(),
4     nn.Linear(128, 128),
5     nn.ReLU(),
6     nn.Linear(5),
7     nn.Softmax(dim=1)
8 )

```

C.

```

1 nn.Sequential(
2     nn.Linear(300, 128),
3     nn.ReLU(),
4     nn.Linear(128, 128),
5     nn.ReLU(),
6     nn.Linear(128, 5),
7     nn.Softmax(dim=1)
8 )

```

D.

```

1 nn.Sequential(
2     nn.Flatten(),
3     nn.Linear(300, 128),
4     nn.ReLU(),
5     nn.Linear(128, 128),
6     nn.ReLU(),
7     nn.Linear(128, 5),
8     nn.Softmax(dim=1)
9 )

```

E. Don't know

**Solution:** We first need to flatten the input image before feeding it into a linear layer. The `Linear` module requires two non-optional arguments: the input size and the output size. The only option that both flattens the input and uses the correct dimensions is option D.

**Question 5**

Consider a 2D convolutional layer with no padding, a stride of 2 (both height and width), no bias, and the kernel

$$K = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}. \quad (3)$$

The input to the layer is

$$X = \begin{bmatrix} 1 & 2 & 0 & 1 \\ 1 & 0 & 1 & 0 \\ 1 & 0 & 2 & 2 \\ 0 & 2 & 1 & 1 \end{bmatrix}. \quad (4)$$

What is the output of the layer?

- A.  $\begin{bmatrix} 3 & 2 \\ 0 & 3 \end{bmatrix}$
- B.  $\begin{bmatrix} 4 & 2 \\ 3 & 6 \end{bmatrix}$
- C.  $\begin{bmatrix} 1 & 0 \\ 3 & 3 \end{bmatrix}$

D.  $\begin{bmatrix} 1 & 3 & 0 \\ 1 & 2 & 3 \\ 3 & 1 & 3 \end{bmatrix}$

E. Don't know

**Solution:** Let the output matrix be  $O$ . With stride 2, the kernel is applied at four positions:

1. Submatrix  $\begin{bmatrix} 1 & 2 \\ 1 & 0 \end{bmatrix}$  gives  $O_{11} = 1 \cdot 1 + 0 \cdot 2 + 0 \cdot 1 + 1 \cdot 0 = 1$ .
2. Submatrix  $\begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}$  gives  $O_{12} = 1 \cdot 0 + 0 \cdot 1 + 0 \cdot 1 + 1 \cdot 0 = 0$ .
3. Submatrix  $\begin{bmatrix} 1 & 0 \\ 0 & 2 \end{bmatrix}$  gives  $O_{21} = 1 \cdot 1 + 0 \cdot 0 + 0 \cdot 0 + 1 \cdot 2 = 3$ .
4. Submatrix  $\begin{bmatrix} 2 & 2 \\ 1 & 1 \end{bmatrix}$  gives  $O_{22} = 1 \cdot 2 + 0 \cdot 2 + 0 \cdot 1 + 1 \cdot 1 = 3$ .

### Question 6

Consider a 2D convolutional layer with padding 1 (height and width), a stride of 2 (height and width), no bias, and 10 kernels of size `(channels_in, kernel_height, kernel_width)=(5, 3, 3)`. For an input tensor of size `(channels, height, width)=(5, 5, 5)`, what is the output size of the layer?

- A. `(channels, height, width) = (10, 2, 2)`
- B. `(channels, height, width) = (10, 3, 3)`
- C. `(channels, height, width) = (10, 5, 5)`
- D. `(channels, height, width) = (5, 3, 3)`
- E. Don't know

**Solution:** The layer has 10 kernels, so the output has 10 channels. Padding 1 turns the  $5 \times 5$  input into  $7 \times 7$ . With a  $3 \times 3$  kernel and a stride of 2, the kernel can start at positions 0, 2, 4 along each dimension. That gives three valid positions horizontally and three vertically. Thus, the output spatial size is  $3 \times 3$ , and the full output size is  $(10, 3, 3)$ .

### Question 7

Consider a fully connected neural network with one hidden layer of size 32 that maps an `input` tensor of size `[B, 10]` to 3 classes. The loss is computed as `loss(model(input), target)`, where `target` is a tensor of class indices of size `[B]`. Which one of the following code snippets correctly defines the network and the loss?

A.

```

1 model = nn.Sequential(
2     nn.Linear(10, 32),
3     nn.ReLU(),
4     nn.Linear(32, 3)
5 )
6 loss = nn.CrossEntropyLoss()

```

B.

```

1 model = nn.Sequential(
2     nn.Linear(10, 32),
3     nn.ReLU(),
4     nn.Linear(32, 3),
5     nn.Softmax(dim=1)
6 )
7 loss = nn.CrossEntropyLoss()

```

C.

```

1 model = nn.Sequential(
2     nn.Linear(10, 32),
3     nn.ReLU(),
4     nn.Linear(32, 3)
5 )
6 loss = nn.Softmax(dim=1)

```

D.

```

1 model = nn.Sequential(
2     nn.Linear(10, 32),
3     nn.ReLU(),
4     nn.Linear(32, 1)
5 )
6 loss = nn.CrossEntropyLoss()

```

E. Don't know

**Solution:** The loss `nn.CrossEntropyLoss()` expects raw logits of shape  $[B, 3]$  and class indices of shape  $[B]$ . Since `CrossEntropyLoss` internally applies `log_softmax`, the network must *not* include a softmax layer. The only option that outputs logits of size  $[B, 3]$  and uses `CrossEntropyLoss()` correctly is option A.

**Question 8**

Consider the (non-scaled) dot-product attention

$$A = \text{softmax}(QK^\top) V,$$

where the softmax function is applied row-wise and given by

$$\text{softmax}(\mathbf{z})_i = \frac{e^{z_i}}{\sum_j e^{z_j}}. \quad (5)$$

Let the query  $Q$ , key  $K$ , and value  $V$  matrices be given by

$$Q = \begin{bmatrix} 1 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}, \quad K = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 0 & 0 \end{bmatrix}, \quad V = \begin{bmatrix} 2 & 0 & 6 & 0 \\ 0 & 4 & 0 & 8 \end{bmatrix}. \quad (6)$$

What is the attention output  $A$ ?

A.  $A = \begin{bmatrix} 1 & 2 & 4 & 6 \\ 0 & 0 & 0 & 0 \end{bmatrix}$

B.  $A = \begin{bmatrix} 2 & 4 & 6 & 8 \\ 0 & 0 & 0 & 0 \end{bmatrix}$

C.  $A = \begin{bmatrix} 1 & 2 & 3 & 4 \\ 1 & 2 & 3 & 4 \end{bmatrix}$

D.  $A = \begin{bmatrix} 2 & 4 & 6 & 8 \\ 2 & 4 & 6 & 8 \end{bmatrix}$

E. Don't know

**Solution:** First, we compute the attention logits

$$QK^\top = \begin{bmatrix} 1 & 1 \\ 0 & 0 \end{bmatrix} \quad (S6)$$

and apply the softmax row-wise

$$\text{softmax}([1 \ 1]) = \begin{bmatrix} \frac{1}{2} & \frac{1}{2} \end{bmatrix} \quad \text{softmax}([0 \ 0]) = \begin{bmatrix} \frac{1}{2} & \frac{1}{2} \end{bmatrix}. \quad (S7)$$

Thus, the attention weights are

$$W = \begin{bmatrix} \frac{1}{2} & \frac{1}{2} \\ \frac{1}{2} & \frac{1}{2} \end{bmatrix}. \quad (S8)$$

Multiply by  $V$  gives us

$$A = \begin{bmatrix} 1 & 2 & 3 & 4 \\ 1 & 2 & 3 & 4 \end{bmatrix}. \quad (S9)$$

**Question 9**

Consider a PyTorch implementation of a fully connected neural network with a single hidden layer of size 16 that takes an input tensor of shape [B, 5]. The network's output is used to parametrise the mean and standard deviation (or a parameter from which it can be computed) of a Gaussian distribution. The loss function computes the negative log probability of the target tensor `target` of shape [B] up to a normalisation constant. Recall that the Gaussian negative log-likelihood can be written as

$$-\log p(t | \mu, \sigma) = \log \sigma + \frac{1}{2} \left( \frac{t - \mu}{\sigma} \right)^2 + \text{const.} \quad (7)$$

Which one of the following code snippets implements this correctly?

A.

```
1 model = nn.Sequential(nn.Linear(5, 16), nn.ReLU(), nn.Linear(16, 1))
2 def loss(y, target):
3     return ((y - target)**2).mean()
```

B.

```
1 model = nn.Sequential(nn.Linear(5,16), nn.ReLU(), nn.Linear(16,2), nn.Softplus())
2 def loss(y, target):
3     return (torch.log(y[:,1]) + 0.5*((target-y[:,0])/y[:,1])**2).mean()
```

C.

```
1 model = nn.Sequential(nn.Linear(5, 16), nn.ReLU(), nn.Linear(16, 2))
2 def loss(y, target):
3     return (torch.log(y[:,1]) + 0.5*((target-y[:,0])/y[:,1])**2).mean()
```

D.

```
1 model = nn.Sequential(nn.Linear(5, 16), nn.ReLU(), nn.Linear(16, 2))
2 def loss(y, target):
3     return (y[:,1] + 0.5*((target - y[:,0])/torch.exp(y[:,1]))**2).mean()
```

E. Don't know

**Solution:** The model must output both a mean and a parameter that guarantees a *positive* standard deviation. Option D outputs two values (mean and log standard deviation), constructs a positive standard deviation using `torch.exp`, and implements the Gaussian negative log-likelihood correctly (up to a constant). Therefore, option D is the correct answer.

**Question 10**

Consider constructing an LSTM for classifying sequences into three classes (a many-to-one mapping). Assume we have one-hot encoded an input sentence to a tensor of shape [T, V], where T is the sequence length and V is the vocabulary size (there is no batch dimension). We want a PyTorch LSTM model that, given such an input, outputs unnormalised logits of shape [3] suitable for use with `nn.CrossEntropyLoss` together with a target tensor containing a single class index of shape [1]. Which one of the following code snippets defines this model correctly?

A.

```
1 class MyLSTM(nn.Module):
2     def __init__(self, V):
3         super().__init__()
4         self.lstm = nn.LSTM(input_size=V, hidden_size=32)
5         self.linear = nn.Linear(32, 3)
6
7     def forward(self, x):
8         x, (h_n, c_n) = self.lstm(x)
9         return self.linear(x[-1])
```

B.

```
1 class MyLSTM(nn.Module):
2     def __init__(self, V):
3         super().__init__()
4         self.lstm = nn.LSTM(input_size=V, hidden_size=32)
5         self.linear = nn.Linear(32, 3)
6
7     def forward(self, x):
8         x, (h_n, c_n) = self.lstm(x)
9         return self.linear(x[:-1])
```

C.

```
1 class MyLSTM(nn.Module):
2     def __init__(self, V):
3         super().__init__()
4         self.lstm = nn.LSTM(input_size=V, hidden_size=32)[-1]
5         self.linear = nn.Linear(32, 3)
6
7     def forward(self, x):
8         x, (h_n, c_n) = self.lstm(x)
9         return self.linear(x)
```

D.

```
1 class MyLSTM(nn.Module):
2     def __init__(self, V):
3         super().__init__()
4         self.lstm = nn.LSTM(input_size=V, hidden_size=32)
5
6     def forward(self, x):
7         x, (h_n, c_n) = self.lstm(x)
8         return x[-1]
```

E. Don't know

**Solution:** For an unbatched input  $x$  of shape  $[T, V]$ , `nn.LSTM(input_size=V, hidden_size=32)(x)` returns as its first output a tensor of shape  $[T, 32]$  (when the input has no batch dimension). For many-to-one sequence classification, we want to use the representation of the *last* time step (shape  $[32]$ ) and map it through a linear layer `Linear(32, 3)` to obtain logits of shape  $[3]$ . Option A does exactly this: it calls the LSTM, takes  $x[-1]$  (last time step, shape  $[32]$ ), and passes it through `self.linear` to get a tensor of shape  $[3]$ .